

Telco Customer Churn Prediction ¶

Intro

The dataset I've chosen to use for this Homework is the Telco Customer Churn dataset from Kaggle.com. This dataset includes a target variable of churn and I will be tuning classification models to see which can predict the best churn of a customer. Churn is if a customer leaves or stays with a company. The business best interest is to retain customers. There dataset includes attribute information of the services a customer is signed up for, customer account information, and demographic information about the customer.

Models used and compared will be K-Nearest Neighbors, Decision Tree, Random Forest, Support Vector Machine (Linear and RBF), Stochastic Gradient Descent, Adaboost, Bagging, Gradient Boosting, ANN, and Stacking.

A few key Metrics to be examined for these will be Accuracy, Recall, R2 and MAE. Summary and details are provided within code below

Import Standard Packages

```
In [1]: #Load packages
import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
from sklearn import preprocessing
import os as os
from sklearn.metrics import mean_squared_error
%matplotlib inline
import sys
from sklearn.metrics import r2_score
from sklearn.metrics import mean_absolute_error
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score
from sklearn.metrics import confusion_matrix, classification_report
from sklearn.model_selection import train_test_split, cross_val_score, KFold
from sklearn.model_selection import GridSearchCV
```

Import Dataset


```
In [2]: Churn=pd.read_csv(r"C:\Users\17857\OneDrive\Desktop\Predictive Analytics\Telco_churn.csv")
```

In [3]: *#view the dataset*
Churn

Out[3]:

| | customerID | gender | SeniorCitizen | Partner | Dependents | tenure | PhoneService | MultipleLin |
|------|------------|--------|---------------|---------|------------|--------|--------------|-------------|
| 0 | 7590-VHVEG | Female | 0 | Yes | No | 1 | No | No phc serv |
| 1 | 5575-GNVDE | Male | 0 | No | No | 34 | Yes | |
| 2 | 3668-QPYBK | Male | 0 | No | No | 2 | Yes | |
| 3 | 7795-CFOCW | Male | 0 | No | No | 45 | No | No phc serv |
| 4 | 9237-HQITU | Female | 0 | No | No | 2 | Yes | |
| ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 7038 | 6840-RESVB | Male | 0 | Yes | Yes | 24 | Yes | Y |
| 7039 | 2234-XADUH | Female | 0 | Yes | Yes | 72 | Yes | Y |
| 7040 | 4801-JAZZL | Female | 0 | Yes | Yes | 11 | No | No phc serv |
| 7041 | 8361-LTMKD | Male | 1 | Yes | No | 4 | Yes | Y |
| 7042 | 3186-AJIEK | Male | 0 | No | No | 66 | Yes | |

7043 rows × 21 columns




In [4]: *#preview data in easier manner to read*
Churn.head()

Out[4]:

| | customerID | gender | SeniorCitizen | Partner | Dependents | tenure | PhoneService | MultipleLines |
|---|------------|--------|---------------|---------|------------|--------|--------------|------------------|
| 0 | 7590-VHVEG | Female | 0 | Yes | No | 1 | No | No phone service |
| 1 | 5575-GNVDE | Male | 0 | No | No | 34 | Yes | No |
| 2 | 3668-QPYBK | Male | 0 | No | No | 2 | Yes | No |
| 3 | 7795-CFOCW | Male | 0 | No | No | 45 | No | No phone service |
| 4 | 9237-HQITU | Female | 0 | No | No | 2 | Yes | No |

5 rows × 21 columns




In [5]: `Churn.tail()`

Out[5]:

| | customerID | gender | SeniorCitizen | Partner | Dependents | tenure | PhoneService | MultipleLin |
|-------------|------------|--------|---------------|---------|------------|--------|--------------|-------------|
| 7038 | 6840-RESVB | Male | 0 | Yes | Yes | 24 | Yes | Y |
| 7039 | 2234-XADUH | Female | 0 | Yes | Yes | 72 | Yes | Y |
| 7040 | 4801-JJAZL | Female | 0 | Yes | Yes | 11 | No | No phc serv |
| 7041 | 8361-LTMKD | Male | 1 | Yes | No | 4 | Yes | Y |
| 7042 | 3186-AJIEK | Male | 0 | No | No | 66 | Yes | |

5 rows × 21 columns



In [6]: `Churn.shape` *#7043 observations and 21 columns*


Out[6]: (7043, 21)

Preprocess the Dataset

In [7]: *#Remove the CustomerID column. It is unnecessary so do not have to deal with it*
#Note: axis=1 denotes that we are referring to a column, not a row
`Churn=Churn.drop('customerID',axis=1)`
`Churn.head(5)`

Out[7]:

| | gender | SeniorCitizen | Partner | Dependents | tenure | PhoneService | MultipleLines | InternetServ |
|----------|--------|---------------|---------|------------|--------|--------------|------------------|--------------|
| 0 | Female | 0 | Yes | No | 1 | No | No phone service | D |
| 1 | Male | 0 | No | No | 34 | Yes | No | D |
| 2 | Male | 0 | No | No | 2 | Yes | No | D |
| 3 | Male | 0 | No | No | 45 | No | No phone service | D |
| 4 | Female | 0 | No | No | 2 | Yes | No | Fiber op |



```
In [8]: #check datatypes
Churn.dtypes
```

```
Out[8]: gender          object
SeniorCitizen      int64
Partner            object
Dependents         object
tenure             int64
PhoneService       object
MultipleLines      object
InternetService    object
OnlineSecurity     object
OnlineBackup       object
DeviceProtection   object
TechSupport        object
StreamingTV        object
StreamingMovies    object
Contract           object
PaperlessBilling   object
PaymentMethod      object
MonthlyCharges     float64
TotalCharges       object
Churn              object
dtype: object
```

```
In [9]: #Senior citizen should be categorical and total charges should be a float variable
Churn['SeniorCitizen'] = Churn['SeniorCitizen'].astype('object')
```

```
In [10]: #After any attempts for many hours total charges would not convert to object so I am going to drop it so I can move on. This would result in too many dummy variable when converting later on
Churn=Churn.drop('TotalCharges',axis=1)
Churn.head(5)
```

Out[10]:

| | gender | SeniorCitizen | Partner | Dependents | tenure | PhoneService | MultipleLines | InternetService |
|---|--------|---------------|---------|------------|--------|--------------|------------------|-----------------|
| 0 | Female | 0 | Yes | No | 1 | No | No phone service | D |
| 1 | Male | 0 | No | No | 34 | Yes | No | D |
| 2 | Male | 0 | No | No | 2 | Yes | No | D |
| 3 | Male | 0 | No | No | 45 | No | No phone service | D |
| 4 | Female | 0 | No | No | 2 | Yes | No | Fiber op |

In [11]: Churn.dtypes

```
Out[11]: gender                object
SeniorCitizen                object
Partner                      object
Dependents                   object
tenure                       int64
PhoneService                 object
MultipleLines                object
InternetService              object
OnlineSecurity               object
OnlineBackup                 object
DeviceProtection             object
TechSupport                  object
StreamingTV                  object
StreamingMovies              object
Contract                     object
PaperlessBilling             object
PaymentMethod                object
MonthlyCharges               float64
Churn                        object
dtype: object
```

Move target variable Churn to front of dataframe

```
In [12]: # designate target variable name & move it to the front of the dataframe
targetName = 'Churn'
#print(targetName)
targetSeries = Churn[targetName]
#print(targetSeries)
#remove target from current location and insert in column number 0
del Churn[targetName]
Churn.insert(0, targetName, targetSeries)
#reprint dataframe and see target is in position 0
Churn.head(5)
```

Out[12]:

| | Churn | gender | SeniorCitizen | Partner | Dependents | tenure | PhoneService | MultipleLines | InternetService |
|---|-------|--------|---------------|---------|------------|--------|--------------|------------------|-----------------|
| 0 | No | Female | 0 | Yes | No | 1 | No | No phone service | |
| 1 | No | Male | 0 | No | No | 34 | Yes | No | |
| 2 | Yes | Male | 0 | No | No | 2 | Yes | No | |
| 3 | No | Male | 0 | No | No | 45 | No | No phone service | |
| 4 | Yes | Female | 0 | No | No | 2 | Yes | No | |

```
In [13]: #Check for missing values  
Churn.isnull().sum() #There are no missing values in the dataset
```

```
Out[13]: Churn          0  
gender          0  
SeniorCitizen   0  
Partner         0  
Dependents      0  
tenure          0  
PhoneService    0  
MultipleLines   0  
InternetService 0  
OnlineSecurity  0  
OnlineBackup    0  
DeviceProtection 0  
TechSupport     0  
StreamingTV     0  
StreamingMovies 0  
Contract        0  
PaperlessBilling 0  
PaymentMethod   0  
MonthlyCharges  0  
dtype: int64
```

Transform the data so factors are dummied

Here We are going to do two preprocessing tasks: 1) change the target column to a numeric since many models in Scikit Learn requires integers/numeric; and 2) create dummy variables for the categorical variables

Change target column to numeric since scikit learn models require it

```
In [14]: # This code turns a character/text target variable into numeric one  
from sklearn import preprocessing  
le_dep = preprocessing.LabelEncoder()  
#to convert into numbers  
Churn['Churn'] = le_dep.fit_transform(Churn['Churn'])
```

Create dummies for categorical variables

```
In [15]: # perform data transformation. Creates dummy variables for categorical variables.

for col in Churn.columns[1:]:
    attName = col
    dType = Churn[col].dtype
    missing = pd.isnull(Churn[col]).any()
    uniqueCount = len(Churn[attName].value_counts(normalize=False))
    #create dummies
    if dType == object:
        Churn = pd.concat([Churn, pd.get_dummies(Churn[col], prefix=col)], axis=1)
    del Churn[attName]
```

```
In [16]: Churn.shape
```

```
Out[16]: (7043, 46)
```


There are now more rows with dummy variables. Up from 20 to 46

```
In [17]: Churn.head()
```

```
Out[17]:
```

| | Churn | tenure | MonthlyCharges | gender_Female | gender_Male | SeniorCitizen_0 | SeniorCitizen_1 |
|---|-------|--------|----------------|---------------|-------------|-----------------|-----------------|
| 0 | 0 | 1 | 29.85 | 1 | 0 | 1 | 0 |
| 1 | 0 | 34 | 56.95 | 0 | 1 | 1 | 0 |
| 2 | 1 | 2 | 53.85 | 0 | 1 | 1 | 0 |
| 3 | 0 | 45 | 42.30 | 0 | 1 | 1 | 0 |
| 4 | 1 | 2 | 70.70 | 1 | 0 | 1 | 0 |

5 rows × 46 columns




Perform some EDA so to get a feel for the data

In [18]: Churn.describe()

Out[18]:

| | Churn | tenure | MonthlyCharges | gender_Female | gender_Male | SeniorCitizen_0 |
|--------------|-------------|-------------|----------------|---------------|-------------|-----------------|
| count | 7043.000000 | 7043.000000 | 7043.000000 | 7043.000000 | 7043.000000 | 7043.000000 |
| mean | 0.265370 | 32.371149 | 64.761692 | 0.495244 | 0.504756 | 0.837853 |
| std | 0.441561 | 24.559481 | 30.090047 | 0.500013 | 0.500013 | 0.368612 |
| min | 0.000000 | 0.000000 | 18.250000 | 0.000000 | 0.000000 | 0.000000 |
| 25% | 0.000000 | 9.000000 | 35.500000 | 0.000000 | 0.000000 | 1.000000 |
| 50% | 0.000000 | 29.000000 | 70.350000 | 0.000000 | 1.000000 | 1.000000 |
| 75% | 1.000000 | 55.000000 | 89.850000 | 1.000000 | 1.000000 | 1.000000 |
| max | 1.000000 | 72.000000 | 118.750000 | 1.000000 | 1.000000 | 1.000000 |

8 rows × 46 columns




The dataset seems to be evenly split between Male and Female. Also Evenly split on a person having a partner or not. 84% of the customers are not Senior citizens so under the age of 65. Over half the customers pay month to month while the mean monthly bill is \$65.

Check the target variable attributes


```
In [19]: #Basic bar chart since the target is binominal  
groupby = Churn.groupby(targetName)  
targetEDA=groupby[targetName].aggregate(len)  
print(targetEDA)  
plt.figure()  
targetEDA.plot(kind='bar', grid=False)  
plt.axhline(0, color='k')
```

```
Churn  
0    5174  
1    1869  
Name: Churn, dtype: int32
```

```
Out[19]: <matplotlib.lines.Line2D at 0x142ce928a88>
```



```
In [20]: # percentage of churn  
1869/5174
```

```
Out[20]: 0.36122922303826827
```

1869 Customers churned meaning they left the company. This is 36% of the observations. 64% of the customers are staying with the company. The ML model should have better than 36% predictive performance

Correlation Matrix

```
In [21]: corr_matrix = Churn.iloc[:,2:].corr()
corr_matrix


#code chunk is attributed as follows:
#Shaw, Chris (2019). How to customize seaborn correlation heat map. Medium.
#Retrieved from https://medium.com/@chrisshaw982/seaborn-correlation-heatmaps-
customized-10246f4f7f4b.

import seaborn as sns

plt.figure(figsize=(46,46)) #need to adjust size as needed.
mask = np.zeros_like(corr_matrix, dtype=np.bool)
mask[np.triu_indices_from(mask)] = True

sns.heatmap(corr_matrix,
            vmin=-1,
            vmax=1,
            cmap='coolwarm',
            annot=True,
            mask=mask)

plt.show()
```



There are a couple correlated variables such as monthly charges along with streaming TVs and Movies. This makes sense as the more services a customer has the higher their bill may be. This is a common theme here the more or less services a customer has correlates with a more positive or negative monthly charge. Also having dependants is correlated with whether or not the customer has a partner as well.

Tuned Modeling with Classification methods

Create training and Test sets using 70/30 split

```
In [22]: # split dataset into testing and training
# column location 1 to end of dataframe are the features.
# column location 0 is the target
features_train, features_test, target_train, target_test = train_test_split(
    Churn.iloc[:,1:].values, Churn.iloc[:,0].values, test_size=0.30, random_state=0)
```

```
In [23]: print(features_test.shape)
print(features_train.shape)
print(target_test.shape)
print(target_train.shape)
```

```
(2113, 45)
(4930, 45)
(2113,)
(4930,)
```

Metrics

Some of the metrics that will be used to evaluate models are as follows

Accuracy score is the most crucial metric in evaluation since it determines how accurate a model is.

R2 is another metric I like to look at. It will show positive and negative correlations in the data

Recall Score shows the true positive positive rate at which the model performs.

MAE will measure the magnitude of errors in the forecast

K-Nearest Neighbor (KNN) Modeling and Tuning

Build and make KNN = 3, 5, and 7

```
In [48]: from sklearn.neighbors import KNeighborsClassifier
from sklearn.preprocessing import StandardScaler
scaler = StandardScaler
knn = KNeighborsClassifier(7) #Making KNN=7
print(knn)
print(scaler)

from sklearn.pipeline import make_pipeline
clf_knn = make_pipeline(StandardScaler(), KNeighborsClassifier(7))

#Train model
clf_knn = clf_knn.fit(features_train, target_train)

#Validate model
target_predicted_knn = clf_knn.predict(features_test)

#Classification Report and Confusion Matrix

print("KNN Accuracy Score", accuracy_score(target_test, target_predicted_knn))
print("R2 Score", r2_score(target_test, target_predicted_knn))
print("Recall Score", recall_score(target_test, target_predicted_knn))
print("Mean Abs Error", mean_absolute_error(target_test, target_predicted_knn))

print(classification_report(target_test, target_predicted_knn))
print(confusion_matrix(target_test, target_predicted_knn))

#extracting true_positives, false_positives, true_negatives, false_negatives
tn, fp, fn, tp = confusion_matrix(target_test, target_predicted_knn).ravel()
print("True Negatives: ",tn)
print("False Positives: ",fp)
print("False Negatives: ",fn)
print("True Positives: ",tp)
```

```

KNeighborsClassifier(algorithm='auto', leaf_size=30, metric='minkowski',
                    metric_params=None, n_jobs=None, n_neighbors=7, p=2,
                    weights='uniform')
<class 'sklearn.preprocessing._data.StandardScaler'>
KNN Accuracy Score 0.7553241836251775
R2 Score -0.2663107989057356
Recall Score 0.5063291139240507
Mean Abs Error 0.24467581637482252

```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.83 | 0.84 | 0.84 | 1560 |
| 1 | 0.53 | 0.51 | 0.52 | 553 |
| accuracy | | | 0.76 | 2113 |
| macro avg | 0.68 | 0.67 | 0.68 | 2113 |
| weighted avg | 0.75 | 0.76 | 0.75 | 2113 |

```

[[1316 244]
 [ 273 280]]
True Negatives: 1316
False Positives: 244
False Negatives: 273
True Positives: 280

```

The accuracy score on this model is 75%

Cross Validate

```

In [25]: #Cross validate KNN with 10-fold cross validation
scores = cross_val_score(clf_knn, features_train, target_train, cv=10)
print("Cross Validation Score for each K", scores)
scores.mean()

```

```

Cross Validation Score for each K [0.76064909 0.77281947 0.76673428 0.7423935
1 0.74239351 0.7525355
0.76064909 0.76267748 0.7505071 0.78701826]

```

```

Out[25]: 0.7598377281947262

```

KNN Results:

I Tuned KNN to 3, 5, and 7 for these three models and the best performance came from the model tuned to 7 KNN with an Accuracy score of 75.5%. This is only slightly better than the other two model with an accuracy score of 73% This model had much better Recall and R2 scores with a lower MAE value.

Decision Tree Modeling and Tuning

Build Decision Tree and make max depth = None, 5, 3

```

In [51]: #decision tree. Call up my model and name it clf
from sklearn import tree
clf_dt = tree.DecisionTreeClassifier()
#Call up the model to see the parameters you can tune (and their default settings)

clf_dt = tree.DecisionTreeClassifier(ccp_alpha=0.0, class_weight=None, criterion='gini',
                                     max_depth=5, max_features=None, max_leaf_nodes=None,
                                     min_impurity_decrease=0.0, min_impurity_split=None,
                                     min_samples_leaf=1, min_samples_split=2,
                                     min_weight_fraction_leaf=0.0, presort='deprecated',
                                     random_state=None, splitter='best')

#Train the model
clf_dt = clf_dt.fit(features_train, target_train)

#Validate
#Predict clf DT model again test data
target_predicted_dt = clf_dt.predict(features_test)

#Classification Report and Confusion Matrix
print("DT Accuracy Score", accuracy_score(target_test, target_predicted_dt))
print("R2 Score", r2_score(target_test, target_predicted_dt))
print("Recall Score", recall_score(target_test, target_predicted_dt))
print("Mean Abs Error", mean_absolute_error(target_test, target_predicted_dt))
print(classification_report(target_test, target_predicted_dt))
print(confusion_matrix(target_test, target_predicted_dt))

tn, fp, fn, tp = confusion_matrix(target_test, target_predicted_dt).ravel()
print("True Negatives: ",tn)
print("False Positives: ",fp)
print("False Negatives: ",fn)
print("True Positives: ",tp)

```

DT Accuracy Score 0.7936583057264552

R2 Score -0.06791394259748684

Recall Score 0.5877034358047016

Mean Abs Error 0.20634169427354473

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.86 | 0.87 | 0.86 | 1560 |
| 1 | 0.61 | 0.59 | 0.60 | 553 |
| accuracy | | | 0.79 | 2113 |
| macro avg | 0.73 | 0.73 | 0.73 | 2113 |
| weighted avg | 0.79 | 0.79 | 0.79 | 2113 |

[[1352 208]

[228 325]]

True Negatives: 1352

False Positives: 208

False Negatives: 228

True Positives: 325

The best models Accuracy score is 79.4%

```
In [27]: #verify DT with 10-fold cross validation
scores = cross_val_score(clf_dt, features_train, target_train, cv=10)
print("Cross Validation Score", scores)
scores.mean()

Cross Validation Score [0.79107505 0.79716024 0.80121704 0.77890467 0.8052738
3 0.78093306
0.76267748 0.77484787 0.79310345 0.8336714 ]

Out[27]: 0.7918864097363083
```

Results:

I Tuned max depth to 3, 5, and None for these three models and the best performance came from the model tuned to 5 max depth with an Accuracy score of 79.2%. This is slightly better than when tuned to 3 it was 78.5% however the default setting was lower at 72%. This model had much better Recall and R2 scores with a lower MAE value.

Random Forest Modeling and Tuning

Build Random Forest making max features auto, 5, and 10

```

In [55]: from sklearn.ensemble import RandomForestClassifier
clf_rf = RandomForestClassifier()

clf_rf = RandomForestClassifier(bootstrap=True, ccp_alpha=0.0, class_weight=None,
                               criterion='gini', max_depth=10, max_features='auto',
                               max_leaf_nodes=None, max_samples=None,
                               min_impurity_decrease=0.0, min_impurity_split=None,
                               min_samples_leaf=1, min_samples_split=2,
                               min_weight_fraction_leaf=0.0, n_estimators=100,
                               n_jobs=None, oob_score=False, random_state=None,
                               verbose=0, warm_start=False)

#Train
clf_rf.fit(features_train, target_train)

#Validate
target_predicted_rf = clf_rf.predict(features_test)

#Classification Report and Confusion Matrix
print("RF Accuracy Score", accuracy_score(target_test, target_predicted_rf))
print("R2 Score", r2_score(target_test, target_predicted_rf))
print("Recall Score", recall_score(target_test, target_predicted_rf))
print("Mean Abs Error", mean_absolute_error(target_test, target_predicted_rf))
print(classification_report(target_test, target_predicted_rf))
print(confusion_matrix(target_test, target_predicted_rf))

tn, fp, fn, tp = confusion_matrix(target_test, target_predicted_rf).ravel()
print("True Negatives: ",tn)
print("False Positives: ",fp)
print("False Negatives: ",fn)
print("True Positives: ",tp)

```

RF Accuracy Score 0.8012304779933743

R2 Score -0.028724440116845162

Recall Score 0.5189873417721519

Mean Abs Error 0.19876952200662565

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.84 | 0.90 | 0.87 | 1560 |
| 1 | 0.65 | 0.52 | 0.58 | 553 |
| accuracy | | | 0.80 | 2113 |
| macro avg | 0.75 | 0.71 | 0.72 | 2113 |
| weighted avg | 0.79 | 0.80 | 0.79 | 2113 |

[[1406 154]

[266 287]]

True Negatives: 1406

False Positives: 154

False Negatives: 266

True Positives: 287

The best models Accuracy score is 79.6%

```
In [29]: #verify DT with 10-fold cross validation
scores = cross_val_score(clf_dt, features_train, target_train, cv=10)
print("Cross Validation Score", scores)
scores.mean()

Cross Validation Score [0.78904665 0.79716024 0.80121704 0.78296146 0.8052738
3 0.78093306
0.76267748 0.77484787 0.79310345 0.8336714 ]

Out[29]: 0.7920892494929006
```

Results:

I tuned max features for the random forest model. As max features increased so did the accuracy of the model. The default setting had an accuracy of 77%. Max features at 5 had an accuracy at 79% and at 10 it was 79.6%

Support Vector Machine (Linear Kernel)

Build SVM using default then tuned to C= 0.01 and C= 0.2

```

In [59]: from sklearn.preprocessing import StandardScaler
from sklearn.pipeline import make_pipeline
from sklearn.svm import SVC
#Build
clf_linearSVC = make_pipeline(StandardScaler(), SVC(kernel='linear', C=0.01, r
andom_state=123))
#print(clf_linearSVC)

#Train
clf_linearSVC.fit(features_train, target_train)

#Validate
target_predicted_linearSVC = clf_linearSVC.predict(features_test)

#Classification Report and Confusion Matrix
print("Accuracy Score", accuracy_score(target_test, target_predicted_linearSVC
))
print("Recall Score", recall_score(target_test, target_predicted_linearSVC))
print("R2 Score", r2_score(target_test, target_predicted_linearSVC))
print("Mean Abs Error", mean_absolute_error(target_test, target_predicted_line
arSVC))

print(classification_report(target_test, target_predicted_linearSVC))
tn, fp, fn, tp = confusion_matrix(target_test, target_predicted_linearSVC).ravel()
print("True Negatives: ",tn)
print("False Positives: ",fp)
print("False Negatives: ",fn)
print("True Positives: ",tp)

```

Accuracy Score 0.8035967818267865

Recall Score 0.5424954792043399

R2 Score -0.01647772059164465

Mean Abs Error 0.19640321817321343

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.85 | 0.90 | 0.87 | 1560 |
| 1 | 0.65 | 0.54 | 0.59 | 553 |
| accuracy | | | 0.80 | 2113 |
| macro avg | 0.75 | 0.72 | 0.73 | 2113 |
| weighted avg | 0.80 | 0.80 | 0.80 | 2113 |

True Negatives: 1398

False Positives: 162

False Negatives: 253

True Positives: 300

The accuracy score for the default model is 69% then 80% then 80.3%

```
In [31]: #verify with 10-fold cross validation
#this will take a long time to run!
scores = cross_val_score(clf_linearSVC, features_train, target_train, cv=10)
print("Cross Validation Score for each K",scores)
scores.mean()
```

Cross Validation Score for each K [0.79918864 0.79716024 0.81541582 0.77890467
0.81744422 0.77890467
0.78904665 0.78904665 0.78701826 0.84989858]

Out[31]: 0.8002028397565922

Tuning the C parameter (soft margin)

```
In [32]: import time
start = time.time()
param_grid={'C': [0.01,0.2]} #trying out two different C values.
clf_linearSVC = SVC(kernel='linear', class_weight='balanced')
grid_svm = GridSearchCV(clf_linearSVC, param_grid,n_jobs=-1, cv=5)
grid_svm.fit(features_train, target_train)
print("SCORES", grid_svm.cv_results_)
print("BEST SCORE", grid_svm.best_score_)
print("BEST PARAM", grid_svm.best_params_)
end = time.time()
print("Time to run", round(end-start), "seconds")
```

SCORES {'mean_fit_time': array([3.18961315, 17.23113465]), 'std_fit_time': array([0.06054432, 2.86264341]), 'mean_score_time': array([0.11840863, 0.06802545]), 'std_score_time': array([0.00749472, 0.01476269]), 'param_C': masked_array(data=[0.01, 0.2],
mask=[False, False],
fill_value='?',
dtype=object), 'params': [{'C': 0.01}, {'C': 0.2}], 'split0_test_score': array([0.73326572, 0.69878296]), 'split1_test_score': array([0.73326572, 0.70182556]), 'split2_test_score': array([0.73022312, 0.69269777]), 'split3_test_score': array([0.72718053, 0.70182556]), 'split4_test_score': array([0.76064909, 0.72109533]), 'mean_test_score': array([0.73691684, 0.70324544]), 'std_test_score': array([0.01207876, 0.00952699]), 'rank_test_score': array([1, 2])}

BEST SCORE 0.736916835699797

BEST PARAM {'C': 0.01}

Time to run 25 seconds

Results

I used the default model which produced an accuracy of 69%. Then as suggested by gridsearch the best c parameter of 0.01 which produced an accuracy of 80.3%. Also the c parameter of 0.2 which was close but produced an accuracy of 80%. The best tuned model was with a C parameter of 0.01

Support Vector Machine (RBF Kernel)

Build using $C = 1$ and $\gamma = 0.01$. Then $C = 5$ and $\gamma = 0.1$. Finally $C = 0.1$ and $\gamma = 0.001$.

```
In [63]: #Build
from sklearn.preprocessing import StandardScaler
from sklearn.pipeline import make_pipeline
from sklearn.svm import SVC
clf_linearSVC = make_pipeline(StandardScaler(), SVC(kernel='rbf', C=1, gamma=
0.01, random_state=123))

#Train
clf_linearSVC.fit(features_train, target_train)

#Validate
target_predicted_linearSVC = clf_linearSVC.predict(features_test)
print("Accuracy Score", accuracy_score(target_test, target_predicted_linearSVC
))
print("Recall Score", recall_score(target_test, target_predicted_linearSVC))
print("R2 Score", r2_score(target_test, target_predicted_linearSVC))
print("Mean Abs Error", mean_absolute_error(target_test, target_predicted_line
arSVC))

print(classification_report(target_test, target_predicted_linearSVC))
tn, fp, fn, tp = confusion_matrix(target_test, target_predicted_linearSVC).ravel()
print("True Negatives: ",tn)
print("False Positives: ",fp)
print("False Negatives: ",fn)
print("True Positives: ",tp)
```

Accuracy Score 0.7969711310932324

Recall Score 0.47920433996383366

R2 Score -0.05076853526220626

Mean Abs Error 0.20302886890676763

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.83 | 0.91 | 0.87 | 1560 |
| 1 | 0.65 | 0.48 | 0.55 | 553 |
| accuracy | | | 0.80 | 2113 |
| macro avg | 0.74 | 0.69 | 0.71 | 2113 |
| weighted avg | 0.78 | 0.80 | 0.79 | 2113 |

True Negatives: 1419

False Positives: 141

False Negatives: 288

True Positives: 265

Results:

The best model by a bit was the original model with $C = 1$ and $\gamma = 0.01$ with an accuracy of 79.7%

Stochastic Gradient Descent

Build tuned models using loss = log, hinge, and modified_huber

```

In [67]: #Build
from sklearn.linear_model import SGDClassifier
from sklearn.pipeline import make_pipeline
import time

start = time.time()
scaler = StandardScaler
clf_sgd_logit = make_pipeline(StandardScaler(), SGDClassifier(loss='log', random_state=123))

#Train
clf_sgd_logit.fit(features_train, target_train)
end = time.time()
print("Training time is", end-start)

#Validate
target_predicted_sgd_logit = clf_sgd_logit.predict(features_test)
print("Accuracy Score", accuracy_score(target_test, target_predicted_sgd_logit))
print("Recall Score", recall_score(target_test, target_predicted_sgd_logit))
print("R2 Score", r2_score(target_test, target_predicted_sgd_logit))
print("Mean Abs Error", mean_absolute_error(target_test, target_predicted_sgd_logit))

print(classification_report(target_test, target_predicted_sgd_logit))
tn, fp, fn, tp = confusion_matrix(target_test, target_predicted_sgd_logit).ravel()
print("True Negatives: ",tn)
print("False Positives: ",fp)
print("False Negatives: ",fn)
print("True Positives: ",tp)

```

Training time is 0.09224843978881836

Accuracy Score 0.7808802650260294

Recall Score 0.5009041591320073

R2 Score -0.13404622803356991

Mean Abs Error 0.21911973497397066

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.83 | 0.88 | 0.86 | 1560 |
| 1 | 0.60 | 0.50 | 0.54 | 553 |
| accuracy | | | 0.78 | 2113 |
| macro avg | 0.71 | 0.69 | 0.70 | 2113 |
| weighted avg | 0.77 | 0.78 | 0.77 | 2113 |

True Negatives: 1373

False Positives: 187

False Negatives: 276

True Positives: 277

Results:

The best model using loss classifiers came from the log model with an accuracy is 78%. The hinge model is 77.2%. And the modified_huber is 74.3%

Adaboost (Use at least two different learners)

Build tuned models using depth = 10 learning rate = 0.01 n estimators = 50. Then max depth = 1 learning rate = 0.1 n estimators = 100. Finally max depth = 1 lr = 0.2 n estimators = 200

```
In [70]: from sklearn.ensemble import AdaBoostClassifier
from sklearn import tree
clf_dt_ab = AdaBoostClassifier(base_estimator=tree.DecisionTreeClassifier(max_
depth=1),
                                algorithm="SAMME.R", n_estimators=200, learning
_rate =0.2, random_state=123)
#Train
clf_dt_ab.fit(features_train, target_train)

#Validate
target_predicted_dt_ab=clf_dt_ab.predict(features_test)
print("Accuracy Score", accuracy_score(target_test, target_predicted_dt_ab))
print("Recall Score", recall_score(target_test, target_predicted_dt_ab))
print("R2 Score", r2_score(target_test, target_predicted_dt_ab))
print("Mean Abs Error", mean_absolute_error(target_test, target_predicted_dt_a
b))

print(classification_report(target_test, target_predicted_dt_ab))
tn, fp, fn, tp = confusion_matrix(target_test, target_predicted_dt_ab).ravel()
print("True Negatives: ",tn)
print("False Positives: ",fp)
print("False Negatives: ",fn)
print("True Positives: ",tp)
```

Accuracy Score 0.8012304779933743

Recall Score 0.5497287522603979

R2 Score -0.028724440116845162

Mean Abs Error 0.19876952200662565

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.85 | 0.89 | 0.87 | 1560 |
| 1 | 0.64 | 0.55 | 0.59 | 553 |
| accuracy | | | 0.80 | 2113 |
| macro avg | 0.74 | 0.72 | 0.73 | 2113 |
| weighted avg | 0.79 | 0.80 | 0.80 | 2113 |

True Negatives: 1389

False Positives: 171

False Negatives: 249

True Positives: 304

Results:

I used the tuned parameters of max depth = 10 learning rate = 0.01 n estimators = 50. Then max depth = 1 learning rate = 0.1 n estimators = 100. Finally max depth = 1 lr = 0.2 n estimators = 200. This final model performed best with an accuracy of 80.1%

Bagging Classifier (choose at least one learner)

Build tuned models using n estimators 100, 500, 1000

```
In [73]: #Build
from sklearn.ensemble import BaggingClassifier
clf_bag = BaggingClassifier(n_estimators=500, random_state=123)

#Train
clf_bag.fit(features_train, target_train)

#Validate
target_predicted_clf_bag=clf_bag.predict(features_test)
print("Accuracy Score", accuracy_score(target_test, target_predicted_clf_bag))
print("Recall Score", recall_score(target_test, target_predicted_clf_bag))
print("R2 Score", r2_score(target_test, target_predicted_clf_bag))
print("Mean Abs Error", mean_absolute_error(target_test, target_predicted_clf_bag))

print(classification_report(target_test, target_predicted_clf_bag))
tn, fp, fn, tp = confusion_matrix(target_test, target_predicted_clf_bag).ravel()
print("True Negatives: ",tn)
print("False Positives: ",fp)
print("False Negatives: ",fn)
print("True Positives: ",tp)
```

Accuracy Score 0.7799337434926644

Recall Score 0.49547920433996384

R2 Score -0.13894491584364999

Mean Abs Error 0.22006625650733555

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.83 | 0.88 | 0.86 | 1560 |
| 1 | 0.60 | 0.50 | 0.54 | 553 |
| accuracy | | | 0.78 | 2113 |
| macro avg | 0.71 | 0.69 | 0.70 | 2113 |
| weighted avg | 0.77 | 0.78 | 0.77 | 2113 |

True Negatives: 1374

False Positives: 186

False Negatives: 279

True Positives: 274

```
In [37]: #verify bagging with cross validation
scores_bag = cross_val_score(clf_bag, features_train, target_train, cv=10, n_jobs=-1)
print("Cross Validation Score for each K",scores_bag)
scores_bag.mean()
```

Cross Validation Score for each K [0.79716024 0.80527383 0.78904665 0.78093306 0.79918864 0.75456389 0.76876268 0.76470588 0.78904665 0.79918864]

Out[37]: 0.7847870182555781

Results

I used the n estimator to tune the model to 100, 500. and 1000. The best model was with n estimator at 500 with an accuracy score of 77.9. But the difference was only in 0.1% each way to the other tuned models.

Gradient Boosting

Using n estimator 100 and learning rate 0.1. Then n estimator 500 and learning rate 0.0. Finally n estimator 100 and learning rate 0.2

```
In [77]: #Build
from sklearn.ensemble import GradientBoostingClassifier
clf_GBC = GradientBoostingClassifier(n_estimators=100, learning_rate=0.1, random_state=123) #default learning rate is 0.1

#Train
clf_GBC.fit(features_train, target_train)

#Validate
target_predicted_GBC=clf_GBC.predict(features_test)
print("Accuracy Score", accuracy_score(target_test, target_predicted_GBC))
print("Recall Score", recall_score(target_test, target_predicted_GBC))
print("R2 Score", r2_score(target_test, target_predicted_GBC))
print("Mean Abs Error", mean_absolute_error(target_test, target_predicted_GBC))

print(classification_report(target_test, target_predicted_GBC))
tn, fp, fn, tp = confusion_matrix(target_test, target_predicted_GBC).ravel()
print("True Negatives: ",tn)
print("False Positives: ",fp)
print("False Negatives: ",fn)
print("True Positives: ",tp)
```

Accuracy Score 0.7936583057264552

Recall Score 0.5298372513562387

R2 Score -0.06791394259748684

Mean Abs Error 0.20634169427354473

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.84 | 0.89 | 0.86 | 1560 |
| 1 | 0.62 | 0.53 | 0.57 | 553 |
| accuracy | | | 0.79 | 2113 |
| macro avg | 0.73 | 0.71 | 0.72 | 2113 |
| weighted avg | 0.79 | 0.79 | 0.79 | 2113 |

True Negatives: 1384

False Positives: 176

False Negatives: 260

True Positives: 293

Results

For tuning this model I first used the n estimator of 100, then n estimator of 500. I saw the accuracy went down so I tuned n estimator back to 100 and adjusted the learning rate from the default of .1 to .2. The best model is the originaly model with an n estimator of 100 and learning rate of 0.1 which results in an accuracy of 79.3%

Extra Trees (Extremely Randomized Trees)

Build using tuned parameter max features default setting, 5, and 10. Then max depth 5, 10. and 3

```
In [80]: #Build
from sklearn.ensemble import ExtraTreesClassifier
clf_xdt = ExtraTreesClassifier(n_estimators= 100, n_jobs=-1, random_state=123,
max_depth=5)

#Train
clf_xdt.fit(features_train, target_train)

#Validate
target_predicted_xdt=clf_xdt.predict(features_test)
print("Accuracy Score", accuracy_score(target_test, target_predicted_xdt))
print("Recall Score", recall_score(target_test, target_predicted_xdt))
print("R2 Score", r2_score(target_test, target_predicted_xdt))
print("Mean Abs Error", mean_absolute_error(target_test, target_predicted_xdt
))

print(classification_report(target_test, target_predicted_xdt))
tn, fp, fn, tp = confusion_matrix(target_test, target_predicted_xdt).ravel()
print("True Negatives: ",tn)
print("False Positives: ",fp)
print("False Negatives: ",fn)
print("True Positives: ",tp)
```

Accuracy Score 0.7927117841930904

Recall Score 0.4213381555153707

R2 Score -0.07281263040756714

Mean Abs Error 0.20728821580690962

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.82 | 0.92 | 0.87 | 1560 |
| 1 | 0.66 | 0.42 | 0.52 | 553 |
| accuracy | | | 0.79 | 2113 |
| macro avg | 0.74 | 0.67 | 0.69 | 2113 |
| weighted avg | 0.78 | 0.79 | 0.78 | 2113 |

True Negatives: 1442

False Positives: 118

False Negatives: 320

True Positives: 233

Results

I began using the tuning parameters of max features using the default setting, 5, then 10. The result were almost identical at 75.6% accuracy score for all these so I decided to see if max depth would make a difference. I did 5, 10, and 3. The best model was tuned to a max depth of 5 with a 79.2% accuracy score

ANN (different hidden layers and nodes)

Build using hidden layer and nodes of 5, 10, and, 3

```
In [83]: #Build
from sklearn.neural_network import MLPClassifier
from sklearn.pipeline import make_pipeline
scaler = StandardScaler

clf_nn = MLPClassifier(hidden_layer_sizes=(5,5), solver="sgd", learning_rate=
"adaptive", max_iter=1000,random_state=0)

#Train
clf_nn.fit(features_train, target_train)

#Validate
target_predicted_clf_nn = clf_nn.predict(features_test)
print("Accuracy Score", accuracy_score(target_test, target_predicted_clf_nn))
print("Recall Score", recall_score(target_test, target_predicted_clf_nn))
print("R2 Score", r2_score(target_test, target_predicted_clf_nn))
print("Mean Abs Error", mean_absolute_error(target_test, target_predicted_clf_
nn))

print(classification_report(target_test, target_predicted_clf_nn))
tn, fp, fn, tp = confusion_matrix(target_test, target_predicted_clf_nn).ravel
()
print("True Negatives: ",tn)
print("False Positives: ",fp)
print("False Negatives: ",fn)
print("True Positives: ",tp)
```

Accuracy Score 0.8045433033601515

Recall Score 0.5171790235081374

R2 Score -0.011579032781564358

Mean Abs Error 0.19545669663984855

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.84 | 0.91 | 0.87 | 1560 |
| 1 | 0.66 | 0.52 | 0.58 | 553 |
| accuracy | | | 0.80 | 2113 |
| macro avg | 0.75 | 0.71 | 0.73 | 2113 |
| weighted avg | 0.79 | 0.80 | 0.80 | 2113 |

True Negatives: 1414

False Positives: 146

False Negatives: 267

True Positives: 286

Results:

For the neural network the best performing model was the one with 5 hidden layers and nodes at 80.4% accuracy rate. I then increased the hidden layers and nodes to 10 but this decreased performance. So finally I decreased to 3 and this also had worse and original performance

Stacking


```
In [41]: from sklearn.linear_model import LogisticRegression
from sklearn.naive_bayes import GaussianNB
from sklearn.ensemble import RandomForestClassifier
from sklearn.ensemble import VotingClassifier
from sklearn.ensemble import AdaBoostClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.preprocessing import StandardScaler
from sklearn.pipeline import make_pipeline

learner_1 = GaussianNB()
learner_2 = DecisionTreeClassifier()
learner_3 = RandomForestClassifier(max_features='auto', n_estimators=100)
learner_4 = AdaBoostClassifier(DecisionTreeClassifier(), algorithm="SAMME.R", n_estimators=100)

stacked_learner = VotingClassifier(estimators=[('lr', learner_1), ('nb', learner_2), ('rf', learner_3), ('adaboost', learner_4)], voting='hard')

for MV, label in zip([learner_1, learner_2, learner_3, learner_4, stacked_learner],
                     ['Naive Bayes', 'Decision Tree', 'Random Forest', 'AdaBoost Decision Tree', 'Second Stage Learner']):
    scores2 = cross_val_score(MV, features_train, target_train, cv=5, scoring='recall')
    scores3 = cross_val_score(MV, features_train, target_train, cv=5, scoring='accuracy')
    scores4 = cross_val_score(MV, features_train, target_train, cv=5, scoring='precision')
    scores5 = cross_val_score(MV, features_train, target_train, cv=5, scoring='r2')
    scores6 = cross_val_score(MV, features_train, target_train, cv=5, scoring='neg_mean_absolute_error')

    print("Recall: %0.2f (+/- %0.2f) [%s]" % (scores2.mean(), scores2.std(), label))
    print("Accuracy: %0.2f (+/- %0.2f) [%s]" % (scores3.mean(), scores3.std(), label))
    print("R2 Score: %0.2f (+/- %0.2f) [%s]" % (scores5.mean(), scores5.std(), label))
    print("MAE: %0.2f (+/- %0.2f) [%s]" % (scores6.mean(), scores6.std(), label))
```

```
Recall: 0.84 (+/- 0.01) [Naive Bayes]
Accuracy: 0.70 (+/- 0.01) [Naive Bayes]
R2 Score: -0.51 (+/- 0.06) [Naive Bayes]
MAE: -0.30 (+/- 0.01) [Naive Bayes]
Recall: 0.50 (+/- 0.02) [Decision Tree]
Accuracy: 0.72 (+/- 0.02) [Decision Tree]
R2 Score: -0.42 (+/- 0.07) [Decision Tree]
MAE: -0.28 (+/- 0.02) [Decision Tree]
Recall: 0.49 (+/- 0.02) [Random Forest]
Accuracy: 0.78 (+/- 0.01) [Random Forest]
R2 Score: -0.12 (+/- 0.09) [Random Forest]
MAE: -0.22 (+/- 0.01) [Random Forest]
Recall: 0.50 (+/- 0.04) [AdaBoost Decision Tree]
Accuracy: 0.76 (+/- 0.02) [AdaBoost Decision Tree]
R2 Score: -0.20 (+/- 0.10) [AdaBoost Decision Tree]
MAE: -0.25 (+/- 0.02) [AdaBoost Decision Tree]
Recall: 0.47 (+/- 0.02) [Second Stage Learner]
Accuracy: 0.77 (+/- 0.01) [Second Stage Learner]
R2 Score: -0.13 (+/- 0.07) [Second Stage Learner]
MAE: -0.23 (+/- 0.01) [Second Stage Learner]
```

Results

When stacking the best stack with an accuracy of 77% is Naive Bayes, Decision Tree, Random Forest, and Adaboost. Some other stacks I experimented with included logistic regression and they and slightly lower accuracy scores performing slightly worse.

Summary

The chart below shows the comparisons of the evaluation metrics of each of the best model for each classifier. All the classifier models seemed to be pretty consistent throughout with this dataset with top accuracy scores being from 78% to 80.5%. When looking at accuracy the best model is the neural network at 80.45%. The model showing the best recall or true positive rate is the decision tree at 58.78%. I thought recall was pretty low for this dataset but it was consistent throughout the classifier methods. The decision had the strongest positive correlation at 6.8% but this is fairly weak. The dataset seems to contain a weak negative to no real correlation at all. This is consistent among the classifiers. The best MAE was the Neural network showing the smallest error.

Overall I think the best model for this dataset is the Neural Network model. It showed the top accuracy which is the main thing I was looking when comparing and also the lowest mean absolute error. It had the third best recall as well.

Some limitation of this dataset is that it looks at a customers current situation and does not look at behaviors for contribution to churn. Such calls with customer service, If they are aware of competitors, satisfaction with service, and other behavioral data may be useful in better predicting customer churn.

| | Accuracy | Recall | R2 | MAE |
|-----------------------------|----------|--------|---------|--------|
| KNN | 0.7553 | 0.5063 | -0.2663 | 0.2447 |
| Decision Tree | 0.7937 | 0.5878 | 0.068 | 0.2063 |
| Random Forest | 0.7989 | 0.5172 | -0.041 | 0.2011 |
| SVM Linear | 0.8036 | 0.5423 | -0.0165 | 0.1964 |
| SVM RBF | 0.797 | 0.4792 | -0.0508 | 0.203 |
| Stochastic Gradient Descent | 0.7809 | 0.5009 | -0.134 | 0.212 |
| Adaboost | 0.8012 | 0.548 | -0.0287 | 0.1988 |
| Bagging | 0.7799 | 0.4955 | -0.1389 | 0.2201 |
| Gradient Boosting | 0.7893 | 0.5244 | -0.0896 | 0.2106 |
| Extra Trees | 0.7927 | 0.4213 | -0.0728 | 0.2073 |
| ANN | 0.8045 | 0.5172 | -0.0116 | 0.1955 |
| Stacking | 0.78 | 0.48 | -0.16 | -0.22 |